


From
Attention
to
Intelligence

TRANSFORMERS

Detailed Notes & Revision Guide

The Core Architecture Behind Modern AI



Attention + Context + Better Representations = Smarter AI

SUBTOPICS COVERED

- ① Transformer Architecture & Components
 - Encoder, Decoder, Transformer block
- ② Multi-Head Attention
 - Why multiple heads, parallel attention, different relationships, combining heads
- ③ Transformer Encoder
 - Encoder block, self-attention, FFN, residual connections, layer norm
- ④ Transformer Decoder
 - Masked self-attention, cross-attention, FFN, residual connections, layer norm
- ⑤ Feed-Forward Network & Normalization
 - Position-wise FFN, activation functions, layer normalization, residual connections
- ⑥ Training Transformers
 - Forward propagation, loss function, backpropagation, optimizers, LR scheduling, teacher forcing
- ⑦ Advanced Transformers & Interview Concepts
 - KV cache, attention complexity, flash attention, MoE, quantization, optimization, interview Q&A
- ⑧ BERT, GPT & Transformer Variants
 - BERT (encoder), GPT (decoder), T5 (encoder-decoder), RoBERTa & others, when to use each
- ⑨ Transformers in LLMs & Generative AI
 - LLMs, text generation, context window, prompting, autoregressive generation, applications
- ⑩ Practical Insights
 - Key differences, advantages, limitations, real-world use cases
- ⑪ Important Formulas & Concepts
 - Attention formula, softmax, positional encoding, complexity (where applicable)
- ⑫ Diagrams & Visuals
 - Architecture diagrams, flow charts, attention mechanism, model comparisons
- ⑬ Quick Revision
 - Key takeaways, do's & don'ts, exam/interview points



Smarter
Applications

Powering the Next Generation of AI



Natural
Language



Gen

SUBTOPIC: Introduction to Transformers1. What are Transformers?Definition:

- Transformers are a type of deep learning model designed to work with sequential data (like text, code, etc.).
- They are based on the self-attention mechanism, which helps the model understand the context of words in a sequence.
- Unlike RNNs, transformers do not process tokens one by one. They can look at all tokens in parallel.

Key idea: Self-Attention

- Allows the model to focus on relevant words in the input sequence.
- Captures long-range dependencies effectively.

Example:

In the sentence "The cat sat on the mat", when predicting "mat", the model can directly look at "cat", "sat", "on", "the" (all tokens) and decide which ones are important.

2. Why Transformers were introduced?Motivation:

- RNNs and LSTMs struggled with long-range dependencies.
- Training was slow due to sequential processing.
- Limited ability to capture global context.
- Transformers solve these issues using self-attention and parallel computation.

Key reasons:


- Better handling of long sequences.
- Faster training (parallelizable).
- Improved performance in NLP tasks.
- Scales well with large data and model sizes.

3. RNN/LSTM LimitationsRNN Limitations:

- Vanishing/exploding gradient problem.
- Struggles with long-term dependencies.
- Sequential processing → slow training.
- Limited memory of past information.

LSTM Limitations:

- Efficient to some extent, but still struggles with long sequences.
- Requires more parameters.
- Slower due to sequential nature.

Transformer Architecture (High Level)Main Components:

- Token Embedding - Converts tokens to vectors.
- Positional Encoding - adds position information.
- Multi-Head Attention - focuses on different parts of the input simultaneously.
- Feed Forward Network - applies non-linearity and learns complex patterns.
- Layer Normalization & Residual Connection - helps with stable training.

★ Remember:

Transformers are more parallelizable and scalable than RNNs, making them ideal for large-scale training.

RNN vs LSTM vs Transformer

Feature	RNN	LSTM	Transformer
Long-term dependencies	Poor	Better	Excellent
Training speed	Slow	Slow	Fast (parallel)
Memory	Limited	Better	Global context
Architecture	Simple	More complex	Attention-based
Scalability	Low	Medium	High

4. Transformer Applications

- Natural Language Processing (NLP) - BERT, GPT, T5, etc.
- Machine Translation - Google Translate, etc.
- Text Summarization - news, articles, documents.
- Question Answering - chatbots, QA systems.
- Text Generation - content creation, code generation.
- Computer Vision - Vision Transformers (ViT).
- Speech Processing - speech-to-text, text-to-speech.
- Multimodal AI - image + text, audio + text, etc.

Real-world examples:

- ChatGPT (GPT series)
- Google BERT (search, NLP)
- DALL-E (image generation)

→ Used in almost every modern AI application today!

SUBTOPIC: Transformer Architecture1. What is Transformer Architecture?


- Transformer architecture is the overall structure of the model, showing how different components (like encoder, decoder, attention, feed forward layers, etc.) work together.
- It defines the flow of information from input to output.
- It is based on the self-attention mechanism instead of recurrent or convolutional layers.

Key idea:

Use self-attention to capture global context and process the entire sequence in parallel.

2. Encoder-Decoder Architecture

- Used in sequence-to-sequence tasks (e.g., translation, summarization, question answering).
- Encoder reads the input sequence and learns a contextual representation.
- Decoder generates the output sequence, using the context from the encoder.

Flow:Example:

Input: "The cat sits"
Output: "Le chat est assis" (French translation)

3. Input → Encoder → Decoder → OutputStep-by-step flow:

- Input tokens are converted to embeddings and positional encodings are added.
- Encoder processes the input using self-attention and feed forward layers.
- Decoder attends to encoder output (cross-attention) and generates the output tokens.
- Output is converted to vocabulary using a linear layer + softmax.

Important:

- ★ Encoder and Decoder have similar structure (stacked layers).
- ★ Decoder uses a masked self-attention layer.
- ★ Positional encoding is needed to retain word order.

4. Main Components① Multi-Head Attention

- Captures different types of relationships.
- Uses multiple attention heads in parallel.

② Feed Forward Network (FFN)


- Two linear layers with ReLU.
- Applied to each position separately.

③ Add & Norm

- Residual connection + Layer Normalization.
- Helps in stable and faster training.

④ Positional Encoding

- Adds information about token position.
- Uses sin/cos functions (or learned embeddings).

5. Original Transformer ArchitectureWhere:

- N = number of layers (e.g., 6 in base model).
- Each layer has 2 sub-layers (attention + FFN).
- Residual connections + LayerNorm after each sub-layer.
- Positional encoding is added to both encoder and decoder inputs.

Key Points:

- ★ Encoder is used to understand the input.
- ★ Decoder generates the output step by step.
- ★ Masked self-attention prevents looking at future tokens.
- ★ Cross-attention helps decoder focus on relevant parts of the input.
- ★ Works well for paraphrase and long-range dependencies.

SUBTOPIC: Tokenization & Input Representation1. Tokens and vocabularyWhat is it?

- Text is broken into smaller units called tokens.
- Each token is a WORD, SUBWORD or character.
- Complete set SET of UNIQUE tokens is called the vocabulary.
- Vocabulary is fixed BEFORE TRAINING and used for converting text to token IDs.

Key Points:

- Tokens can be words, subwords or characters.
- Vocabulary size is usually 10K - 100K+.
- Special tokens like [PAD], [UNK], [CLS], [SEP] are also part of the vocabulary.

Example (word level):

"The cat is on the mat" → [The] [cat] [is] [on] [the] [mat]

Vocabulary = {The, cat, is, on, the, mat, ...}

REMEMBER:

Remembers: Larger token too flexibility but require more memory and computation.

2. Word/subword tokenizationWhat is it?

- Word tokening: it the text into words.
- Subword tokenization (e.g. BPE, WordPiece) splits rare words into smaller pieces.
- Helps handle unknown words and reduces vocabulary size.

Common methods:

- Word-level (simple, but large vocab).
- Subword-level (BPE, WordPiece, SentencePiece).
- Character-level (rarely used in modern ML LLMs).

Method	Out of vocab?	Vocab size	Example
Word-level	High	Large	"unbelievable"
BPE	Low	Medium	"un ## believable"
WordPiece	Low	Medium	"un ## believable"
Character-level	Very low	Very large	"u n b e l i e v a b l e"

EX (BPE / WordPiece):

[[##Bel] [##iev] [##able]

tokenization is most commonly in most modern LLMs.

3. Token IDsWhat is it?

- Each token from the vocabulary is mapped to a unique integer ID.
- These IDs are used as input to the model.
- Special tokens have fixed IDs (e.g. [PAD]=0, [UNK]=1).

Example:

Token	ID
[PAD]	0
[UNK]	1
The	2
cat	3
is	4
...	...

Key Points:

- Token IDs are integers.
- Range: 0 to vocab_size - 1.
- Used for embedding lookup.
- Padding token helps in batching sequences.

4. EmbeddingsWhat is it?

- Embeddings convert token IDs into dense vectors.
- Capture semantic meaning and context.
- Learned during training.

Key Points:

- Embedding dimension = 768 (BERT base), 1024 (BERT large)
- Tokens with similar meaning have similar embeddings.
- Can be initialized randomly or with pre-trained vectors.

Token IDs

2
3
4
...

EMBEDDING
LAYER

Embedding Vector (d=768)

[0.12, -0.03, 0.45, ...]
[0.08, 0.21, -0.11, ...]
[-0.05, 0.32, 0.17, ...]
...

5. Input EmbeddingsWhat is it?

- Combines token embeddings with positional information.
- Helps the model understand the order of tokens.
- Input to the first encoder layer.

Formula:

Input Embedding = Token Embedding + Positional Encoding

Token IDs

2
3
4
...

Token
Embedding
(d = 768)

Positional
Encoding
(d = 768)

Input
Embedding
(d = 768)

REMEMBER:

- Positional encoding can be sinusoidal or learned.
- Final input shape: (batch_size, seq_len, d_model).

SUBTOPIC: Positional Encoding1. Why position information is neededWhat is the problem?

- Transformer uses self-attention, which is permutation invariant (doesn't consider the order of tokens).
- Without position information, the model cannot distinguish between "dog bites man" and "man bites dog".
- Hence, we need to inject the position of each token into the input representation.

Key idea:

Add positional information to the input embeddings so the model knows the relative/absolute position of each token in the sequence.

2. Sinusoidal positional encodingHow it works?

- Uses sine and cosine functions of different frequencies.
- Allows the model to learn relative positions easily.

$$PE(pos, 2i) = \sin\left(\frac{pos}{10000^{2i/d_{model}}}\right)$$

$$PE(pos, 2i + 1) = \cos\left(\frac{pos}{10000^{2i/d_{model}}}\right)$$

pos = position of the token
i = dimension
d_{model} = model dimension

Properties:

- Produces unique encoding for each position.
- Allows the model to easily learn relative positions (using linear operations).
- Works well for longer sequences (extrapolates).

Example (d_{model} = 6):

Position	Dim 0 (sin)	Dim 1 (cos)	Dim 2 (sin)	Dim 3 (cos)	Dim 4 (sin)	Dim 5 (cos)
0	0.000	1.000	0.000	1.000	0.000	1.000
1	0.841	0.540	0.046	0.999	0.003	1.000
2	0.909	-0.416	0.093	0.996	0.006	1.000
3	0.141	-0.990	0.139	0.990	0.009	1.000
...	...	...	...	...	...	...

(values rounded for simplicity)

3. Learned positional embeddingsHow it works?

- Position is represented using a learnable embedding vector (like word embeddings).
- A separate embedding is maintained for every position (up to a maximum length).

Formula:

$$X = \text{Token Embeddings} + \text{Position Embeddings}$$


(Position embeddings are learned parameters)

Advantages:

- Can learn task-specific position representations.
- Often gives better performance in practice.

Disadvantages:

- Doesn't generalize well to longer sequences (limited to max length seen during training).
- Uses more parameters.

Example:4. Position information in Transformers

- Positional encodings (sinusoidal or learned) are added to the input embeddings at the start.
- The combined representation is passed to the encoder and decoder.
- This allows the model to use position information while computing attention.

Key points:

- Both encoder and decoder use positional encoding.
- In some models, relative positional encoding is used (e.g., Transformer-XL, T5).
- Position information helps with understanding order, structure and long-range dependencies.

SUBTOPIC: Self-Attention Mechanism1. What is attention?

- Attention is a mechanism that allows the model to focus on relevant parts of the input sequence while processing a token.
- It helps in capturing relationships between different tokens, no matter how far apart they are.
- In Transformers, self-attention is used, where the attention is computed within the same sequence (not from an external source).

★ Key idea:

Instead of processing tokens one by one (like RNNs), attention lets the model look at all tokens in parallel and decide which ones are important.

2. Query (Q)

- Query represents the current token (or position) for which we want to find relevant information.
- It is compared with all keys to calculate attention scores.
- Shape: (seq_len, d_model)

Example:

For the word "cat", its query vector is a learned representation that is used to find relevant keys.

3. Key (K)

- Key represents the other tokens in the sequence.
- It is used to compare with the query.
- Each token has its own key vector.
- Shape: (seq_len, d_model)

Example:

All tokens (e.g. "The", "cat", "is", "...") have key vectors. The query is matched with these keys to get attention scores.

4. Value (V)

- Value contains the actual information of each token.
- After computing attention scores, the values are combined to get the final representation.
- Shape: (seq_len, d_model)

Example:

value vector holds the information to the next layer.

5. Attention calculationStep 1: Compute Q, K, V

- Apply 3 separate linear transformations to the input embeddings (X).

$$Q = XW_Q \quad K = XW_K \quad V = XW_V$$

where W_Q, W_K, W_V are learnable weight matrices.

Input: X (seq_len × d_model)

Output: Q, K, V (seq_len × d_model)

Step 2: Compute attention scores

- Use scaled dot-product to get similarity between query and key.

$$\text{Attention scores} = \frac{QK^T}{\sqrt{d_k}} \quad (d_k = \text{dimension of keys/queries})$$

Step 3: Apply softmax


- Convert scores into probabilities (so they sum to 1).

$$\text{Attention weights} = \text{softmax} \left(\frac{QK^T}{\sqrt{d_k}} \right)$$

Step 4: Get weighted sum of values

- Multiply attention weights with the value vectors.

$$\text{Output} = \text{Attention weights} \times V$$

Simple example (1 token attending d.t.e. to 4 tokens):

Each token can attend to all tokens in the sequence and get a context-aware representation.

Key points:

- ★ Q finds relevant tokens.
- ★ K is used for comparison.
- ★ V contains actual information.
- ★ Scaling ($1/\sqrt{d_k}$) prevents large dot products.
- ★ Softmax converts scores to probabilities.
- ★ Self-attention is computed in parallel for all tokens.

SUBTOPIC: Scaled Dot-Product Attention

1. What is Scaled Dot-Product Attention?

Definition:

- Scaled dot-product attention is a mechanism used in Transformers to compute attention scores between tokens in a sequence.
- It helps the model decide which tokens are more relevant while processing a token.
- Uses dot product of query and key, with scaling and softmax to get attention weights.

Key idea:

Compute similarity between query and keys, scale it, apply softmax and use the weights to get a weighted sum of values.

2. Q, K, V Matrices

What are they?

- Q (Query):** Represents what the current token is looking for.
- K (Key):** Represents what each token contains.
- V (Value):** Represents the actual information to be aggregated.

Shapes:

- If input sequence length = n , model dimension = d_{model}
- $Q, K, V \in \mathbb{R}^{n \times d_{\text{model}}}$

Example:

For a sentence with 4 tokens and $d_{\text{model}} = 64$, Q, K, V each have shape 4×64 .

3. Dot Product

What is it?

- Dot product measures similarity between two vectors.
- Higher value $\rightarrow$ more similar.

Formula:

$$a \cdot b = \sum_{i=1}^d a_i b_i$$

In attention:

- We compute $Q \cdot K^T$ (query with all keys).
- ... matrix of shape $n \times n$.

Key points:

- Each token attends to other tokens.
- ... higher attention score (before scaling).

4. Scaling

Why scaling?

- Dot products can get very large as d_{model} increases.
- This pushes softmax into regions with very small gradients.
- To avoid this, we scale the dot product by $\sqrt{d_k}$.

Formula:

$$\text{Scaled Dot Product} = \frac{Q \cdot K^T}{\sqrt{d_k}}$$

Where d_k is the dimension of key vectors (usually d_{model}/h).

Key point:

Scaling keeps the values in a range where softmax works properly and gradients are stable.

5. Softmax

What is it?

- Converts the scaled scores into probabilities.
- Ensures the weights sum to 1 and are between 0 and 1.

Formula:

$$\text{softmax}(x_i) = \frac{e^{x_i}}{\sum_{j=1}^n e^{x_j}}$$

Why softmax?

- Normalizes the scores.
- Gives more focus to relevant tokens.

Properties:

- Output values are in $(0, 1)$.
- Sum of all values = 1.
- Differentiable (important for training).

6. Attention Score Formula

Complete formula:

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{Q \cdot K^T}{\sqrt{d_k}}\right) V$$

Steps:


- Compute dot product: $Q \cdot K^T$
- Scale by $\sqrt{d_k}$
- Apply softmax to get attention weights
- Multiply weights with V
- Get final output (weighted sum)

Example (simplified):

$$Q = [1, 0] \quad K = [1, 0] \quad V = [2, 3] \quad d_k = 2$$

- $Q \cdot K^T = 1$
- Scaled = $1/\sqrt{2} = 0.707$
- $\text{softmax}(0.707) = 1$ (single token example)
- Output = $1 \times V = [2, 3]$

Flow (single head):



Transformers - Page 7

SUBTOPIC: Multi-Head Attention

Insta ID: @tech_1_k

1. Why multiple attention heads?

What is the idea?

- A single attention head has limited capacity to focus on different types of relationships in the input.
- Using multiple heads allows the model to attend to information from different representation subspaces at the same time.
- Each head can learn different patterns (e.g. syntactic, semantic, positional etc.).

Key points:


- Improves model's ability to capture diverse features.
- Helps in learning richer representations.
- Leads to better performance in many NLP tasks.

2. Parallel attention

What is it?

- All attention heads are computed in parallel (not one after another).
- The input is projected into multiple heads simultaneously.
- This makes the computation efficient and faster.

Parallel computation:



Key points:

- Heads are computed in parallel using matrix operations.
- Uses parallelizable GPU/TPU computation.
- Reduces training time compared to sequential attention.

3. Different relationships captured by each head

What does it mean?


- Each head learns to focus on different types of relationships between tokens.
- Some heads may capture syntactic relations (e.g. subject-verb).
- Some heads may capture semantic relations (e.g. coreference).
- Some heads may focus on positional or local patterns.

Examples:

- Head 1: subject-verb agreement.
- Head 2: coreference (e.g. "it" → "dog").
- Head 3: local context (nearby words).
- Head 4: long-range dependencies.

Illustration:

Sentence: "The cat that chased the mouse was hungry."



Key points:

- Each head has its own set of learned parameters.
- Heads work on different subspaces of the input.
- Combined, they give a richer understanding of the context.

4. Combining attention heads

How it works?

1. Compute attention for each head in parallel.
2. Concatenate the outputs from all heads.
3. Apply a linear transformation to get the final output.

Formula:


$$\text{MultiHead}(Q, K, V) = \text{Concat}(\text{head}_1, \text{head}_2, \dots, \text{head}_h) W^O$$

$$\text{head}_i = \text{Attention}(QW_i^Q, KW_i^K, VW_i^V)$$

Where:

- n = number of heads
- W_i^Q, W_i^K, W_i^V = separate projection matrices for each head
- W^O = output projection matrix

Architecture:



Key points:

- Concatenation combines all head outputs.
- Linear layer (W^O) projects back to model dimension.
- Helps the model integrate information from all heads.

Quick takeaway:

- ★ Multi-head attention = multiple parallel + concatenation + linear layer
→ Captures diverse relationships and improves performance

SUBTOPIC: Transformer Encoder1. What is Transformer Encoder?


- The transformer encoder is the part of the Transformer model that reads the input sequence and converts it into a sequence of contextualized representations.
- It uses self-attention and feed-forward layers to understand the meaning of each token in the context of the entire sequence.
- The encoder is composed of multiple identical layers (stacked one after another).

Key idea:

The encoder captures the context from the input sequence and produces rich representations (embeddings) for each token.

2. Encoder Block

- The encoder is built using a repeating block called the encoder layer.
- Each layer has 2 main sub-layers:
 - Multi-Head Self-Attention
 - Feed-Forward Network (MLP)
- Both sub-layers are wrapped with residual connections and layer normalization.

Encoder Layer (Block)3. Self-Attention

- Self-attention allows each token to attend to all tokens in the input sequence.
- It computes attention scores to find how much focus each token should give to other tokens.
- It helps in capturing long-range dependencies and contextual meaning.

Key points:

- Each token attends to all tokens (including itself).
- Attention weights are computed using Q, K, V.
- The encoder uses multiple attention heads in parallel.

4. Feed-Forward Network

- After self-attention, each token is passed through a feed-forward network (MLP) independently.
- It consists of two linear transformations with a ReLU activation in between.

$$\text{FFN}(x) = \max(0, xW_1 + b_1)W_2 + b_2$$

Where:


- $W_1 \in \mathbb{R}^{d_{\text{model}} \times d_{\text{ff}}}$
- $W_2 \in \mathbb{R}^{d_{\text{ff}} \times d_{\text{model}}}$
- d_{ff} is usually $4 \times d_{\text{model}}$ (e.g. 2048 if $d_{\text{model}} = 512$).

5. Residual Connections

- Residual connections add the input of each sub-layer to its output.
- It helps in better gradient flow and avoids the vanishing gradient problem.

Formula:

$$\text{Output} = \text{Layer}(x) + x$$

6. Layer Normalization


- Layer normalization normalizes the activations of each layer across the features.
- It stabilizes training and improves convergence.

Formula (for a single token):

$$\text{LN}(x) = \frac{x - \mu}{\sqrt{\sigma^2 + \epsilon}} \gamma + \beta$$

Where:

- μ = mean
- σ^2 = variance
- γ, β = learnable parameters
- ϵ = small constant

7. Encoder Layer FlowKey points:

- The encoder has N identical layers (e.g. 6 layers).
- Each layer maintains the same dimension (d_{model}).
- Position encoding is added to input embeddings (before first layer).
- The sequence is converted into contextualized representations.

1. Transformer Decoder

What is it?

- The decoder is the part of the Transformer model that generates the output sequence (e.g. translation, text generation) one token at a time.
- It uses masked self-attention, cross-attention and feed-forward layers.
- Each decoder block is stacked multiple times (N times) to form the full decoder.

Key idea:


The decoder attends to both the already generated tokens and the encoder's output to produce the next token.

2. Decoder Block

What is it?

- The decoder block is the building block of the decoder.
- It has 3 main sub-layers:
 1. Masked Self-Attention
 2. Cross-Attention
 3. Feed-Forward Network
- Each sub-layer is followed by a residual connection and layer normalization.

Decoder Block Architecture:



- The output of one block is passed to the next decoder block.
- Stacked N times.

Key points:

- Uses masked self-Attention (to prevent looking ahead).
- Uses cross-attention (to attend to encoder output).
- Helps generate tokens in a left-to-right manner.

3. Masked Self-Attention

What is it?

- Similar to regular self-attention, but with a mask.
- The mask prevents the model from attending to future tokens.
- Ensures autoregressive (left-to-right) generation.

Masking:

Uses a triangular <BOS> matrix where future positions are set to $-\infty$ (or a very large negative value).

6. Autoregressive Generation

What is it?

tokens one by one, in a sequential way. based on the previously generated tokens. to maintain causality.

token.
using the decoder.
sequence and repeat.

Example (mask matrix for seq len = 4):

	1	2	3	4
1	0	$-\infty$	$-\infty$	$-\infty$
2	0	0	$-\infty$	$-\infty$
3	0	0	0	$-\infty$
4	0	0	0	0

$-\infty$ = very large negative value (ignored in softmax)

Key points:

- Allows each position to attend only to previous positions (including itself).
- Maintains the order of the sequence.

4. Cross-Attention

What is it?

- Used to attend to the encoder's output (source sequence).
- Helps the decoder focus on relevant parts of the input.
- The queries come from the decoder, while keys and values come from the encoder output.

Formula:

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V$$

Where:

- Q = decoder's queries
- K, V = encoder's keys and values
- d_k = dimension of keys/queries

Key points:

- Connects decoder with encoder output.
- Helps in understanding the input context.

5. Feed-Forward Network

What is it?

- Same as in the encoder.
- Applies two linear transformations with a ReLU in between.
- Works independently on each position.

Formula:

$$\text{FFN}(x) = \max(0, xW_1 + b_1)W_2 + b_2$$

Where:

- $W_1 \in \mathbb{R}^{d_{\text{model}} \times d_{\text{ff}}}$
- $W_2 \in \mathbb{R}^{d_{\text{ff}} \times d_{\text{model}}}$
- d_{ff} is usually $4 \times d_{\text{model}}$ (e.g. 2048 if $d_{\text{model}} = 512$).

Example:

Input: <BOS> I love NLP

Output (step by step):

<BOS> $\rightarrow$ I $\rightarrow$ love $\rightarrow$ NLP $\rightarrow$ <EOS>

Key points:

- Ensures proper sequence generation.
- Used in tasks like summarization, machine translation, etc.
- Often combined with beam search or greedy decoding.

SUBTOPIC: Feed-Forward Neural Network & Normalization1. Position-wise Feed-Forward Network (FFN)What is it?


- The FFN is a simple 2-layer fully connected network applied independently to each position (token).
- It increases the model's capacity and introduces non-linearity.
- Same weights are used across all positions.

Formula:

$$\text{FFN}(x) = \max(0, xW_1 + b_1) W_2 + b_2$$

Where:

- $W_1 \in \mathbb{R}^{d_{\text{model}} \times d_{\text{ff}}}$, $W_2 \in \mathbb{R}^{d_{\text{ff}} \times d_{\text{model}}}$
- d_{ff} is usually $4 \times d_{\text{model}}$ (e.g. 2048 if $d_{\text{model}} = 512$).

Architecture:2. Activation FunctionsCommon functions:

Function	Formula	Properties
ReLU	$f(x) = \max(0, x)$	• Simple • Fast • Helps with vanishing gradient
GELU	$f(x) = x\Phi(x)$ (Φ = Gaussian CDF)	• Smoother than ReLU • Better performance in Many NLP tasks
Tanh	$f(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}}$	• Output in $(-1, 1)$ • Zero centered

★ GELU is commonly used in modern transformer models (e.g. BERT, GPT).

3. Layer NormalizationWhat is it?

- Normalizes the activations across the feature dimension for each token (instead of across the batch).
- Helps stabilize training and makes the model less sensitive to learning rate.

Formula:

$$\text{LN}(x) = \frac{x - \mu}{\sqrt{\sigma^2 + \epsilon}} \gamma + \beta$$

Where:

- μ = mean (across features)
- σ^2 = variance (across features)
- γ, β = learnable parameters
- ϵ = small constant

Key takeaways:

- LN and FFN (with residual connection).
- Faster training.
- Works well with small batch sizes.

4. Residual ConnectionsWhat is it?


- Adds the input of a sub-layer to its output.
- Helps in better gradient flow and prevents vanishing gradient problem.
- Allows training of deeper models.

Formula:


$$\text{Output} = \text{Layer}(x) + x$$

Where:

- $\text{Layer}(x)$ = output of sub-layer (e.g. attention or FFN)
- x = input to the sub-layer

Example (simplified):5. Transformer Block FlowHow it all fits together?

1. Multi-Head Attention
2. Add & Norm (residual connection + layer norm)
3. Feed-Forward Network (position-wise)
4. Add & Norm (residual connection + layer norm)

Transformer BlockKey takeaways:

- FFN adds non-linearity and capacity.
- Activation functions control the learning behavior.
- LayerNorm stabilizes training (feature-wise).
- Residual connections help gradient flow.
- Together, they make the Transformer

SUBTOPIC: Training Transformers1. Forward PropagationWhat is it?

- Forward propagation is the process of passing the input through the model to get the output.
- In a Transformer, input tokens are converted to embeddings, then passed through multiple layers (encoder/decoder).
- The final layer produces logits, which are converted to probabilities using softmax.

Flow:Key points:

- Each layer consists of self-attention, add & norm, and feed-forward network.
- During training, we compute the loss using the output probabilities and target tokens.

2. Loss FunctionWhat is it?

- The loss function measures how well the model's predictions match the actual target values.
- For Language modeling, we use Cross-Entropy Loss.

Formula:

$$L = -\frac{1}{N} \sum_{i=1}^N \sum_{v=1}^V y_{i,v} \log(p_{i,v})$$

Where:

- N = number of tokens in the batch
- V = vocabulary size
- $y_{i,v} = 1$ if token v is the target, else 0
- $p_{i,v}$ = predicted probability for token v

Key points:

- Lower loss means better predictions.
- In practice, we use label smoothing to prevent overfitting.

3. BackpropagationWhat is it?

- Backpropagation calculates gradients of the loss w.r.t. model parameters using the chain rule.
- It helps update the weights to reduce the loss.
- In Transformers, gradients flow through attention, FFN, and embedding layers.

Key steps:

1. Compute loss and output gradients.
2. Backpropagate through linear layers, attention and FFN. using optimizer.

derivatives.
gradients for large models.

4. OptimizersWhat are they?

- Optimizers update model parameters based on gradients.
- Common optimizers: Adam, AdamW, SGD.

Adam update rule:

$$m_t = \beta_1 m_{t-1} + (1 - \beta_1) g_t$$

$$v_t = \beta_2 v_{t-1} + (1 - \beta_2) g_t^2$$

$$\theta_t = \theta_{t-1} - \alpha \frac{\hat{m}_t}{\sqrt{\hat{v}_t} + \epsilon}$$

Where:


- g_t = gradient
- m_t, v_t = moment estimates
- β_1, β_2 = decay rates
- α = learning rate
- ϵ = small constant

Key points:

- Adam is widely used for training Transformers.
- AdamW decouples weight decay for better generalization.

5. Learning-Rate SchedulingWhat is it?

- Adjusts the learning rate during training.
- Helps the model converge faster and more stably.
- Common strategies: warmup + decay, linear decay, cosine decay.

Example: Warmup + Cosine DecayFormula (cosine decay):

$$lr = lr_{min} + \frac{1}{2} (lr_{max} - lr_{min}) \times \left(1 + \cos\left(\frac{\pi t}{T}\right)\right)$$

Where:


- t = current step
- T = total steps
- lr_{max}, lr_{min} = max/min lr

Key points:

- Warmup prevents instability at the start.
- Cosine decay gives better final performance.

6. Teacher ForcingWhat is it?

- During training, we feed the ground truth token as input to the decoder (instead of the model's prediction).
- Helps the model learn faster and reduces error accumulation.

With Teacher Forcing:Key points:

- Used only during training.
- During inference, we use the model's own predictions.
- Important for stable and faster training.

★ Remember:


Teacher forcing is a training technique, not used during inference.

SUBTOPIC: Pre-training & Fine-Tuning1. Pre-trainingWhat is it?

- Pre-training is the process of training a large model on a massive amount of unlabeled data.
- The model learns general language patterns, world knowledge and useful representations.
- It is done once and provides a strong starting point for many downstream tasks.

Key points:


- Uses large-scale, diverse data (e.g. web text).
- Trains with self-supervised objectives (e.g. masked LM).
- Results in a general-purpose model (e.g. BERT, GPT).

Example architecture:2. Fine-tuningWhat is it?

- Fine-tuning uses the pre-trained model to a specific downstream task using labeled data.
- The model's weights are updated (usually with a smaller learning rate).
- It helps the model perform well on the target task without training from scratch.

Key points:

- Uses task-specific labeled data.
- Only a few layers or all layers are updated.
- Faster and better performance than training from scratch.

Example:3. Transfer LearningWhat is it?

- Transfer learning uses knowledge gained from one task (source) to improve performance on another task (target).
- In transformers, this is often done using pre-trained models and fine-tuning on a new task.

Key points:


- Target data is limited.
- Reduces computational cost.
- Improves performance.

4. Self-Supervised LearningWhat is it?

- Self-supervised learning learns from the data itself without explicit labels.
- The model creates its own training signals (pseudo-labels).
- Common methods: masked language modeling, next sentence prediction, contrastive learning.

Key points:

- Uses unlabeled data.
- Learns useful representations.
- Forms the basis for most modern LLMs.

Example: Masked Language Modeling (MLM)5. Task-Specific ModelsWhat is it?

- Task-specific models are pre-trained models fine-tuned for a particular task or domain.
- They may also be trained from scratch when enough labeled data is available.

Key points:

- Use task-specific data and objectives.
- Can be specialized for different domains (e.g. medical, legal).
- Sometimes use parameter-efficient tuning (e.g. LoRA, adapters).

Example:


Model	Type	Task
BERT	Encoder	Classification, QA
GPT	Decoder	Generation, Chat
T5	Encoder-Decoder	Translation, Summarization

★ Important:

- Pre-training gives general knowledge.
- Fine-tuning adapts to a specific task.
- Transfer learning helps when data is limited.
- Self-supervised learning creates training signals without labels.
- Task-specific models are optimized for particular tasks.

SUBTOPIC: BERT, GPT & Transformer Variants1. BERT - Encoder-onlyWhat is it?

- BERT (Bidirectional Encoder Representations from Transformers) is an encoder-only model.
- It understands context from both left and right (bidirectional) using masked language modeling (MLM).
- Pre-trained on large text corpora and fine-tuned for downstream tasks.

Architecture (simplified)Key points:

- Encoder-only (bidirectional attention).
- Uses MLM during pre-training.
- Great for understanding context and classification tasks.
- Example: BERT-base, BERT-large.

2. GPT - Decoder-onlyWhat is it?


- GPT (Generative Pre-trained Transformer) is a decoder-only model.
- It uses causal (masked) self-attention to look only at previous tokens.
- Trained with next token prediction (autoregressive).

Architecture (simplified)Key points:

- Decoder-only (causal attention).
- Uses next token prediction during pre-training.
- Best for text generation and creative tasks.
- Example: GPT-2, GPT-3, GPT-4, GPT-4o.

3. T5 - Encoder-DecoderWhat is it?

- T5 (Text-to-Text Transfer Transformer) uses an encoder-decoder architecture.
- It frames all tasks (like translation, summarization, QA) as text-to-text problems.

Architecture (simplified)ints:

- Architecture.
- in bias (instead of absolute).
- framework for multiple tasks.
- T5-large.

4. RoBERTa and other variantsWhat is it?

- RoBERTa is an improved version of BERT.
- Key changes: removed next sentence prediction (NSP), used dynamic masking, larger batch sizes, more training steps and data.
- Other variants include ALBERT, DistilBERT, ELECTRA, DeBERTa, etc. Each brings improvements in efficiency, training or performance.

Comparison of popular variants:

Model	Main idea / Improvement	Architecture
RoBERTa	No NSP, dynamic masking, larger data/batch size	Encoder-only
ALBERT	Parameter sharing, factorized embeddings	Encoder-only
DistilBERT	Knowledge distillation, smaller model	Encoder-only
ELECTRA	Replaced token detection	Encoder-only
DeBERTa	Disentangled attention, better positional encoding	Encoder-only

Key points:

- Different training objectives and optimizations.
- Aim to improve performance, efficiency or both.
- Still mostly encoder-only (except some hybrid models).

5. When to use each architecture?

Model	When to use?
BERT	Text understanding, classification, QA, sentiment analysis, document analysis.
GPT	Text generation, creative writing, summarization, chatbots, assistance, code generation.
T5	General purpose (translation, summarization, QA, classification, etc.) using text-to-text.
RoBERTa Variants	When you need better performance/efficiency for understanding tasks (e.g. classification, QA) with large datasets.

Quick Comparison


- BERT → Encoder-only → Understanding
- GPT → Decoder-only → Generation
- T5 → Encoder-Decoder → Text-to-Text
- RoBERTa → Encoder-only → Improved BERT

★ Remember:

The best model depends on your task, and performance vs. efficiency trade-off.

SUBTOPIC: Transformers in LLMs & Generative AI1. Large Language Models (LLMs)What is it?


- LLMs are transformer-based models trained on massive text data (internet books, code, etc.).
- They learn language patterns, world knowledge and can generate human-like text.
- Examples: GPT-3/4, LLaMA, PaLM, Claude, etc.

Architecture (simplified)Key points:

- Trained with self-supervised learning (next token prediction).
- Very large number of parameters (billions).
- Learns general language understanding & generation.

2. Text GenerationWhat is it?

- Text generation is the process of generating coherent, contextually relevant text from a given prompt.
- Uses the learned patterns from training data.
- Can be used for chat, summarization, code generation, etc.

Basic process:Key points:

- Works one token at a time (autoregressive).
- Quality depends on prompt and model size.
- Can be controlled using parameters (temperature, top-k, etc.).

3. Context WindowWhat is it?

- Context window is the maximum number of tokens the model can consider at once.
- Includes both input prompt + generated tokens.
- Larger Context window = better long context understanding.

Typical sizes:

Model	Context window (approx.)
GPT-3.5	4K tokens
GPT-4	8K - 128K tokens
Claude	100K+ tokens
	8K - 128K tokens

Important:

If the context exceeds the limit, older tokens may be truncated or lost (depending on the model's strategy).

with better reasoning and consistency.
using longer input windows or techniques
(Augmented Generation).

4. PromptingWhat is it?

- Prompting is the technique of giving instructions or examples to the model to get the desired output.
- Works well with few-shot or zero-shot examples.
- Can be used for tasks like summarization, translation, Q&A, etc.

Types of prompting:


1. Zero-shot → no examples
2. Few-shot → few examples
3. Chain-of-thought → step-by-step reasoning
4. System prompt → set behavior/persona

Key points:

- Clear and specific prompts give better results.
- Can control tone, format, style, etc.
- Helps in reducing hallucinations (to some extent).

5. Autoregressive GenerationWhat is it?


- The model generates text one token at a time.
- Each new token is predicted based on the previous tokens.
- Uses causal (masked) attention to avoid looking ahead.

Key points:

- Ensures coherent and contextually relevant output.
- Used in most modern LLMs (GPT, LLaMA, etc.).
- Can be influenced by sampling techniques (temperature, top-k, top-p).

6. LLM ApplicationsCommon use cases:


- Chatbots / virtual assistants
- Content generation
- Summarization
- Translation
- Code generation
- Question answering
- Search & RAG
- Creative writing

Example: ChatbotKey takeaways:

- LLMs + transformers = powerful text generation.
- Context window limits how much the model
- Prompting helps guide the model's behavior
- Autoregressive generation enables coherence
- Widely used in real-world applications like

SUBTOPIC: Advanced Transformers & Interview Concepts1. KV CacheWhat is it?

- KV cache stores the key and value tensors from previous tokens during inference.
- Reuses them for future tokens instead of recomputing.
- Reduces computation and speeds up generation.

Architecture (simplified):Key points:

- Mainly used in autoregressive generation.
- Saves memory and compute.
- For long sequences, cache size grows ($\propto$ sequence length).

2. Attention ComplexityWhat is it?

- Self-attention has quadratic complexity w.r.t. sequence length.
- Each token attends to all other tokens.

Complexity:

Time: $O(n^2 \cdot d)$ Space: $O(n^2)$

(n = sequence length, d = hidden size)

Optimizations:

- Flash Attention (reduce memory access).
- Sparse attention (attend to subset).
- Linear attention (approximation).

Comparison:

Method	Time Complexity	Space Complexity
Standard Attention	$O(n^2)$	$O(n^2)$
Flash Attention	$O(n^2)$	$O(n)$
Linear Attention	$O(n)$	$O(n)$

Key point:

- Attention is the main computational bottleneck in long sequences.

3. Long-Context TransformersWhat is it?

- Models that can handle very long input sequences (e.g. 32K, 128K).
- Requires architectural and training changes.
- Helps in tasks like document understanding, code analysis, etc.

Techniques:

- RoPE / ALiBi (better positional encoding).
- RoPE for extrapolation).

videos, code, etc.

reasoning.

Example:

GPT-4o $\rightarrow$ 128K
Claude $\rightarrow$ 200K+

4. Flash AttentionWhat is it?


- An optimized attention algorithm.
- Reduces memory usage and improves speed.
- Avoids storing the full attention matrix.

How it works?

1. Computes attention in blocks (tiling).
2. Uses efficient GPU memory access.
3. Recomputes values instead of storing all intermediate results.

Key points:

- 2-4x faster than standard attention.
- Significantly less memory usage.
- Widely used in modern LLMs (e.g. Llama 3, Mistral).

5. Mixture of Experts (MoE)What is it?

- Uses multiple expert networks instead of one large model.
- A gating network selects relevant experts for each token.
- Improves efficiency and scalability.

Architecture:Key idea:

Only a few experts are activated per token.

Benefits:

- Better scalability.
- Higher performance with fewer active parameters.
- Used in models like Switch Transformer (Google).

6. QuantizationWhat is it?

- Reduces model size and memory usage by using lower precision.
- E.g. FP32 $\rightarrow$ INT8 / INT4 / FP16.
- Minimal impact on performance (with proper calibration).

Precision	Size Reduction	Typical Use
FP32	1x	Training / High accuracy
FP16	2x	Mixed precision training
INT8	4x	Inference (common)
INT4	8x	Edge / Mobile / Faster inference

Techniques:

- Post-training quantization (PTQ)
- Quantization aware training (QAT)
- GPTQ / AWQ

Key point:

- Quantization enables deployment on resource-constrained devices.

7. Transformer OptimizationTechniques:

1. Mixed precision training (FP16 / BF16).
2. Gradient accumulation (larger batch size).
3. Gradient checkpointing (save memory).
4. Better optimizers (AdamW, Lion).
5. Efficient attention (Flash Attention, sparse).

Remember:

- Focus on the concept, not just memorizing facts.
- Be ready to explain trade-offs and real-world use cases.
- Draw diagrams when needed.

8. Common Transformer Interview Questions

1. Why do we use positional encoding in Transformers?
2. What is the difference between encoder-only, decoder-only and encoder-decoder?
3. How does self-attention work? Explain with an example.
4. What is the role of layer normalization and residual connections?
5. How does Flash Attention improve performance?
6. What are MoE models? Why are they useful?
7. How does quantization affect model performance?
8. Compare BERT, GPT and T5. When would you use each?
9. How do long-context models handle very long inputs?
10. What are the trade-offs in model optimization?